COHESITY

Version 4.3.0
March 2020

Cohesity Imanis Data Release Notes

Table of Contents

Imanis Data software supports the following in RELEASE 4.3.0:

APPLICATION	SUPPORTED ENTERPRISE VERSIONS	
	CLOUDERA	HORTONWORKS
HADOOP	6.2.0	3.1.0
	6.1.1	3.0.1
	6.0.1	2.6.5
	5.16.1	2.6.4
	5.14.1	2.6.0
	5.12.2	2.5.4
	5.8.2	2.5.0
	5.7.6	2.3.6
	-	2.2.0

APPLICATION	SUPPORTED ENTERPRISE VERSIONS	
	APACHE	DATASTAX ENTERPRISES
CASSANDRA	3.11	6.7**

APPLICATION	SUPPORTED ENTERPRISE VERSIONS	
	APACHE	DATASTAX ENTERPRISES
	3.10	5.1.10
	3.0	5.1.9
	2.1.11	5.1.8
	2.1.8	5.1.7
	2.1	5.1.2
	-	5.1.0
	-	5.0
	-	4.8

** Refer to the Cassandra limitations section for more information.

APPLICATION	SUPPORTED ENTERPRISE VERSIONS
	COUCHBASE INC.
COUCHBASE	6.0
	5.5.2
	5.1.0
	5.0.1

APPLICATION	SUPPORTED ENTERPRISE VERSIONS
	COUCHBASE INC.
	4.6.3
	4.6.2
	4.5.0
	4.1.1

APPLICATION	SUPPORTED ENTERPRISE VERSIONS
	MONGODB INC.
MONGODB	4.2**
	4.0
	3.6
	3.4

** Refer to the MongoDB limitations section for limitations of MongoDB 4.2 support.

1 New Features

- Support for MongoDB version 4.2

2 Known Issues/ Limitations

This section describes the known issues or limitations in the current release.

All Applications

- While editing a job or workflow, the value in the Suffix field cannot be changed. This restriction is conceived by design and applicable across all the applications where ever the suffix field is enabled
- HTTPS feature for GUI is supported only for RHEL 7.x and CentOS 7.x
- Supported kernel versions are 3.10.0-862.el7.x86_64 and 3.10.0-1062.12.1.el7.x86_64. Any other kernel version between these versions is not supported. Example : 3.10.0-957.el7.x86_64 Is not supported.

Hadoop HDFS

- DLM workflow is not supported
- Support for backup and restore on ACLs and other extended attributes is not available in the current release
- SOLR feature is not supported in HDFS Direct Replication workflow. The feature is available in Mirroring, Pipeline, and Recovery
- In case of Imanis Data upgrade from releases prior to 4.0.0, edit existing Hadoop data repositories and specify HBase principal in the following format: hbase.superuser@your-realm.com. This step is required only for setups which have Kerberized Hadoop repositories and use HBase application. Not required otherwise.
- Direct Replication is not supported for TDE.
- SOLR is not supported for TDE.

Hive

- User Attributes are not preserved after data sampling-masking.
- Hive Direct Replication is not supported for TDE.
- Restore jobs with suffix and overwrite flag, will overwrite any existing data on target cluster.
- Data migration across versions is not supported. For example, data recovery from 5.8.2 to 5.12.2 is not supported.

Hbase

- File system attributes are not replicated during data recovery
- Data migration across versions is not supported. For example, data recovery from 5.8.2 to 5.12.2 is not supported. In case you upgrade to a new version of Hbase, it is required to recreate all the workflows that are associated with the data repository
- HBase Direct Replication is not supported for TDE.

- Data migration across versions is not supported. For example, data recovery from 5.8.2 to 5.12.2 is not supported.

Cassandra

- User must re-discover and re-verify a data repository when configuration settings of the data repository are modified. For example, if changes are made to cassandra.yaml, cassandra-env.sh, or dse.yaml file of a data repository, then the data repository must be re-discovered and re-verified.
- If a Cassandra node is decommissioned and Imanis data compaction is performed, then Imanis Data takes a full backup of the data from the nodes.
- In the case of Point-in-Time (PIT) recovery failure, user must clean up the failed table directories on the primary nodes before rerunning PIT recovery
- Point-in-Time (PIT) recovery is not supported if a table or keyspace is deleted from the source data repository
- In case of PIT restore of views, user needs to run nodetool flush <keyspace_name.viewname> on all nodes after restoring data and before restarting Cassandra process. This step must be performed for all views that are restored
- For DataStax Point-in-Time (PIT) recovery limitations, refer to the following: https://docs.datastax.com/en/opscenter/6.1/opsc/online_help/services/opscBackupServicePointInTimeRestore.html
- Backup is not supported for the following system keyspaces:
  - system
  - system_auth
  - system_distributed
  - system_schema
  - system_traces
  - dse_leases
  - dse_perf
  - dse_security
  - dse_system
  - solr_admin

  Restore of system keyspaces such as OpsCenter is only supported for the same cluster from which it was backed up. To enable backup for any of the preceding keyspaces, remove the particular keyspace from the list of system keyspaces defined in the following configuration file: $INSTALL_DIR/conf/cassandra-conf.xml

- For data migration across Cassandra clusters with different versions, schema changes in source cluster are not supported during the lifetime of the Data migration workflow
- Semi-colon (';') is not supported in column names

- Compaction will skip tables that are encrypted or contains DSE data types like line, point, and so on
- Tables containing UDTs in Cassandra version 3.0 (from 3.0.0 to 3.0.9) and 3.1 to 3.8 are not supported in Recovery Sandbox feature
- Data recovery from Cassandra 2.x to 3.x is not supported if the backed up tables contain dropped columns.
- DSE Tiered Storage support does not work in DSE 5.1 Release. This is a limitation in DSE code.
- For Cassandra 2.x and DSE 4.x, backup recovery is not supported for tables where all the columns are a part of the Primary Key.
- In case of Graphs, jobtag restore fails if suffix is used. This is because individual graph keyspaces cannot be renamed. This issue occurs during jobtag recovery process when global suffix is applied to all keyspaces thus wrongly renaming the '_system' keyspace for a graph. For example, we need to backup two types of keyspaces for a graph. If the name of the graph is ‘exampleGraph’, then we should backup ‘exampleGraph’ and ‘exampleGraph_system’. However, during jobtag recovery, if global suffix is applied to all keyspaces it wrongly renames the keyspaces 'exampleGraph_suffix' and 'exampleGraph_system_suffix' for a graph
- Support for non-default store_type, algorithm, and protocol in cassandra.yaml is not available in this release.
- **For DataStax Enterprise 6.7, the following limitations are applicable:
  - Some of the new features of DSE 6.7 such as NodeSync are not supported.
  - Backup jobs from DSE 5.x cannot be recovered to DSE versions 6.0.x to 6.0.6.
    Recovery to versions before and after the above mentioned versions work correctly.
  - Cassandra agent will not delete temporary directories on the primary nodes if the job gets killed or the primary node goes down during restore. In such cases, the user has to manually delete the temporary directories on the primary nodes.
    The directories are created in each of the storage directories on the primary nodes and are named by the restore job's UUID.

Couchbase

- If a bucket is deleted from the source cluster and a new bucket is created with the same name, then backup of that particular bucket is not supported.
  For example, when a bucket is deleted from Couchbase source cluster, the bucket is still present in the source list of Imanis Data backup job. If a new bucket with the same name is created, the bucket is backed up automatically. As a user, you must edit the job to remove the bucket from the source list or not create a bucket with the same name as the deleted bucket.
- Filtering, Sampling, Masking Options feature is not supported in Point-in-Time (PIT) Recovery
- Point-in-time (PIT) recovery is not supported if Imanis Data cluster is Kerberized
- If bucket name contains special character %, then the backup of that particular bucket is not supported
- Masks are only applied at the leaf nodes of a JSON document. Intermediate nodes of a JSON document will not be masked. For example, given the following JSON documents (DocA and DocB), a mask

specified on the “credit_card_num” attribute will result in only DocA getting masked, where “credit_card_num” is a leaf node. Doc2, on the other hand, where “credit_card_num” is an intermediate node with its children as “type” (Amex, Visa, Mastercard, etc.) and “Number” will not be masked. Imanis Data software supports backup and recovery of Couchbase type buckets. Backup and Recovery of ephemeral and memcached buckets is not supported.

MongoDB

- Truncate collection is not supported
- MongoDB supports a special type of collections called 'Capped Collections'. https://docs.mongodb.com/manual/core/capped-collections/. Imanis MongoDB Connector has limited support for capped collections. A full backup (and recovery in case of mirror and pipeline workflows) is performed for capped collections in every run.
- Backup and recovery of MongoDB Views is not supported
- Indexes with the Unique constraint are not restored on the target collection.
- If a Shard Key has unique constraint set on it, then restore will not apply the unique constraint on the target collection.
- Standalone MongoDB clusters are not supported.
- Distributed transactions (introduced in MongoDB 4.2) are not supported. https://docs.mongodb.com/manual/core/transactions/#transactions-and-atomicity
- Shared Key modifications (introduced in MongoDB 4.2) are not supported. https://docs.mongodb.com/manual/core/sharding-shard-key/#update-shard-key

Your Feedback

Was this document helpful? Send us your feedback!

ABOUT COHESITY

Cohesity makes your data work for you by consolidating secondary storage silos onto a hyperconverged, web-scale data platform that spans both private and public clouds. Enterprise customers begin by radically streamlining their backup and data protection, then converge file and object services, test/dev instances, and analytic functions to provide a global data store. Cohesity counts many Global 1000 companies and federal agencies among its rapidly growing customer base and was named to Forbes’

"Next Billion-Dollar Startups 2017," LinkedIn's "Startups: The 50 Industry Disruptors You Need to Know Now," and CRN's "2017 Emerging Vendors in Storage" lists.

For more information, visit our website and blog, follow us on Twitter and LinkedIn and like us on Facebook.